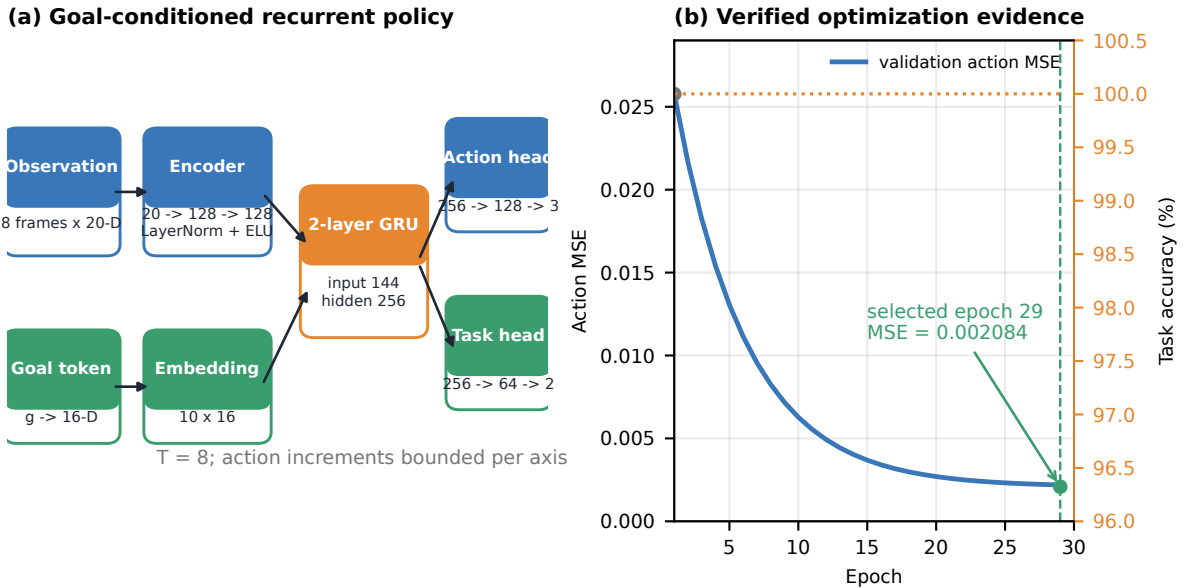
(a) Goal-conditioned recurrent policy
Observation
8 frames x 20-D
Encoder
20 -> 128 -> 128
LayerNorm + ELU
2-layer GRU
input 144
hidden 256
Action head
256 -> 128 -> 3
Goal token
g -> 16-D
Embedding
10 x 16
Task head
256 -> 64 -> 2
T = 8; action increments bounded per axis
(b) Verified optimization evidence
validation action MSE
selected epoch 29
MSE = 0.002084
Action MSE
Task accuracy (%)
Epoch
0.025
0.020
0.015
0.010
0.005
0.000
100.5
100.0
99.5
99.0
98.5
98.0
97.5
97.0
96.5
96.0
5
10
15
20
25
30